

# Lecture 2: Standard library

---

Alexander Stanovoy

February 22, 2022

[alex.stanovoy@gmail.com](mailto:alex.stanovoy@gmail.com)

## In this lecture

- `Option` and `Result`
- `Vec` and `VecDeque`
- `BTreeMap` and `BTreeSet`
- `HashMap` and `HashSet`
- `BinaryHeap`
- `LinkedList`
- `String` and `&str`
- `Box` and `Rc`

# Option and Result

# Option<sup>1</sup> and Result<sup>2</sup>

Let's remember their definitions:

```
enum Option<T> { // 'Maybe' from functional languages
    Some(T),
    None,
}
```

```
enum Result<T, E> {
    Ok(T),
    Err(E),
}
```

---

<sup>1</sup>[Option documentation](#)

<sup>2</sup>[Result documentation](#)

Matching Option:

```
let result = Some("string");  
match result {  
    Some(s) => println!("String inside: {s}"),  
    None => println!("Ooops, no value"),  
}
```

## Option API

Useful functions `.unwrap()` and `.expect()`:

```
fn unwrap(self) -> T;  
fn expect(self, msg: &str) -> T;
```

Useful functions `.unwrap()` and `.expect()`:

```
let opt = Some(22022022);
assert!(opt.is_some());
assert!(!opt.is_none());
assert_eq!(opt.unwrap(), 22022022);
let x = opt.unwrap(); // Copy!

let newest_opt: Option<i32> = None;
// newest_opt.expect("I'll panic!");

let new_opt = Some(Vec::<i32>::new());
assert_eq!(new_opt.unwrap(), Vec::<i32>::new());
// error[E0382]: use of moved value: `new_opt`
// let x = new_opt.unwrap(); // Clone!
```

# Option API

We have a magic function:

```
fn as_ref(&self) -> Option<&T>; // &self is &Option<T>
```

Let's solve a problem:

```
let new_opt = Some(Vec::<i32>::new());
assert_eq!(new_opt.unwrap(), Vec::<i32>::new());
// error[E0382]: use of moved value: `new_opt`
// let x = new_opt.unwrap(); // Clone!

let opt_ref = Some(Vec::<i32>::new());
assert_eq!(new_opt.as_ref().unwrap(), &Vec::<i32>::new());
let x = new_opt.unwrap(); // We used reference!
// There's also .as_mut() function
```

That means if type implements `Copy`, `Option` also implements `Copy`.



## Option API

We can map `Option<T>` to `Option<U>`:

```
fn map<U, F>(self, f: F) -> Option<U>;
```

Example:

```
let maybe_some_string = Some(String::from("Hello, World!"));  
// `Option::map` takes self *by value*,  
// consuming `maybe_some_string`  
let maybe_some_len = maybe_some_string.map(|s| s.len());  
assert_eq!(maybe_some_len, Some(13));
```

## Option API

There's **A LOT** of different `Option` functions, enabling us to write beautiful functional code:

```
fn map_or<U, F>(self, default: U, f: F) -> U;
fn map_or_else<U, D, F>(self, default: D, f: F) -> U;
fn unwrap_or(self, default: T) -> T;
fn unwrap_or_else<F>(self, f: F) -> T;
fn and<U>(self, optb: Option<U>) -> Option<U>;
fn and_then<U, F>(self, f: F) -> Option<U>;
fn or(self, optb: Option<T>) -> Option<T>;
fn or_else<F>(self, f: F) -> Option<T>;
fn xor(self, optb: Option<T>) -> Option<T>;
fn zip<U>(self, other: Option<U>) -> Option<(T, U)>;
```

It's recommended for you to study the documentation and try to avoid `match` where possible.

## Option and ownership

There's two cool methods to control ownership of the value inside:

```
fn take(&mut self) -> Option<T>;  
fn replace(&mut self, value: T) -> Option<T>;  
fn insert(&mut self, value: T) -> &mut T;
```

The first one takes the value out of the **Option**, leaving a **None** in its place.

The second one replaces the value inside with the given one, returning **Option** of the old value.

The third one inserts a value into the **Option**, then returns a mutable reference to it.

## Option API and ownership: take

```
struct Node<T> {  
    elem: T,  
    next: Option<Box<Node<T>>>,  
}  
  
pub struct List<T> {  
    head: Option<Box<Node<T>>>,  
}  
  
impl<T> List<T> {  
    pub fn pop(&mut self) -> Option<T> {  
        self.head.take().map(|node| {  
            self.head = node.next;  
            node.elem  
        })  
    }  
}
```

## Option and optimizations

Rust guarantees to optimize the following types `T` such that `Option<T>` has the same size as `T`:

- `Box<T>`
- `&T`
- `&mut T`
- `fn, extern "C" fn`
- `#[repr(transparent)]` struct around one of the types in this list.
- `num::NonZero*`
- `ptr::NonNull<T>`

This is called the “null pointer optimization” or NPO.

# Result

Functions return `Result` whenever errors are expected and recoverable. In the `std` crate, `Result` is most prominently used for I/O.

**Results must be used!** A common problem with using return values to indicate errors is that it is easy to ignore the return value, thus failing to handle the error. `Result` is annotated with the `#[must_use]` attribute, which will cause the compiler to issue a warning when a `Result` value is ignored.<sup>3</sup>

---

<sup>3</sup>The Error Model

We can match it as a regular enum:

```
let version = Ok("1.1.14");  
match version {  
    Ok(v) => println!("working with version: {:?}", v),  
    Err(e) => println!("error: version empty"),  
}
```

## Result API

We have pretty the same functionality as in `Option`:

```
fn is_ok(&self) -> bool;
fn is_err(&self) -> bool;
fn unwrap(self) -> T;
fn unwrap_err(self) -> E;
fn expect_err(self, msg: &str) -> E;
fn expect(self, msg: &str) -> T;
fn as_ref(&self) -> Result<&T, &E>;
fn as_mut(&mut self) -> Result<&mut T, &mut E>;
fn map<U, F>(self, op: F) -> Result<U, E>;
fn map_err<F, O>(self, op: O) -> Result<T, F>;
// And so on
```

It's recommended for you to study the documentation and try to avoid `match` where possible.



## Operator ?

Consider the following structure:

```
struct Info {  
    name: String,  
    age: i32,  
}
```

## Operator ?

```
fn write_info(info: &Info) -> io::Result<()> {  
    let mut file = match File::create("my_best_friends.txt") {  
        Err(e) => return Err(e),  
        Ok(f) => f,  
    };  
    if let Err(e) = file  
        .write_all(format!("name: {}\n", info.name)  
        .as_bytes()) {  
        return Err(e)  
    }  
    if let Err(e) = file  
        .write_all(format!("age: {}\n", info.age)  
        .as_bytes()) {  
        return Err(e)  
    }  
    Ok(())  
}
```

## Operator ?

We can use the ? operator to make the code smaller!

```
fn write_info(info: &Info) -> io::Result<()> {  
    let mut file = File::create("my_best_friends.txt")?;  
    file.write_all(format!("name: {}\n", info.name).as_bytes())?;  
    file.write_all(format!("age: {}\n", info.age).as_bytes())?;  
    Ok(())  
}
```

Beautiful, isn't it?

We can use it for **Option** too!

## {Result, Option}::transpose

It's time to link Result and Option together:

```
// self is Option<Result<T, E>>
fn transpose(self) -> Result<Option<T>, E>;

// self is Result<Option<T>, E>
fn transpose(self) -> Option<Result<T, E>>;
```

## {Result, Option}::transpose

```
fn read_until_empty() -> io::Result<String> {  
    let mut input = stdin().lines();  
    let mut output = String::new();  
  
    // input.next() gives Option<Result<String>>  
    while let Some(line) = input.next() {  
        let line = line?;  
        if line.is_empty() {  
            break;  
        }  
        output.push_str(&line);  
    }  
  
    Ok(output)  
}
```

## {Result, Option}::transpose

```
fn read_until_empty() -> io::Result<String> {  
    let mut input = stdin().lines();  
    let mut output = String::new();  
  
    // input.next() gives Option<Result<String>>  
    while let Some(line) = input.next().transpose()? {  
        if line.is_empty() {  
            break;  
        }  
        output.push_str(&line);  
    }  
  
    Ok(output)  
}
```

Finally: containers

## About Rust containers in general

Rust containers have some special properties:



## About Rust containers in general

Rust containers have some special properties:

- We don't want to allocate till we need this allocation. Allocations are costly, and usually, it is the last thing we want to do.

## About Rust containers in general

Rust containers have some special properties:

- We don't want to allocate till we need this allocation. Allocations are costly, and usually, it is the last thing we want to do.
- All safe methods try to return **Option** or **Result** where possible, thus containers are predictable and safe.

## About Rust containers in general

Rust containers have some special properties:

- We don't want to allocate till we need this allocation. Allocations are costly, and usually, it is the last thing we want to do.
- All safe methods try to return **Option** or **Result** where possible, thus containers are predictable and safe.
- Some methods **panic!** when some bad things happen (failed allocation, out-of-bounds in “operator [ ]”).

## About Rust containers in general

Rust containers have some special properties:

- We don't want to allocate till we need this allocation. Allocations are costly, and usually, it is the last thing we want to do.
- All safe methods try to return **Option** or **Result** where possible, thus containers are predictable and safe.
- Some methods **panic!** when some bad things happen (failed allocation, out-of-bounds in “operator [ ]”).
- We don't have iterators like in C++, but we have references to elements.

## About Rust containers in general

Rust containers have some special properties:

- We don't want to allocate till we need this allocation. Allocations are costly, and usually, it is the last thing we want to do.
- All safe methods try to return **Option** or **Result** where possible, thus containers are predictable and safe.
- Some methods **panic!** when some bad things happen (failed allocation, out-of-bounds in “operator [ ]”).
- We don't have iterators like in C++, but we have references to elements.

These differences will **affect what algorithms and data structures we can use in the standard library** of Rust.

## Fallible allocations

Interesting fact: since we panic on fallible allocations, we can't use Rust in the Linux kernel currently. But the work being done.

- [RFC 2116](#)
- [Email to Linus Torvalds about that](#)

Just implementation of default dynamic array. The same as in C++, not including some Rust-related differences. Asymptotics are the same as you might expect.

---

<sup>4</sup>Danila Kutenin blog (Russian). [On pqsort, New sorting algorithm for LLVM libcxx](#)

<sup>5</sup>[GoingNative 2013 C++ Seasoning](#)

<sup>6</sup>[Vec documentation](#)

Just implementation of default dynamic array. The same as in C++, not including some Rust-related differences. Asymptotics are the same as you might expect.

Some basic algorithms. **They work on slices too:**

---

<sup>4</sup>Danila Kutenin blog (Russian). [On pqsort, New sorting algorithm for LLVM libcxx](#)

<sup>5</sup>[GoingNative 2013 C++ Seasoning](#)

<sup>6</sup>[Vec documentation](#)



Just implementation of default dynamic array. The same as in C++, not including some Rust-related differences. Asymptotics are the same as you might expect.

Some basic algorithms. **They work on slices too:**

- **sort**, uses modified timsort,  $O(N \log N)$ .

---

<sup>4</sup>Danila Kutenin blog (Russian). [On pqsort, New sorting algorithm for LLVM libcxx](#)

<sup>5</sup>[GoingNative 2013 C++ Seasoning](#)

<sup>6</sup>[Vec documentation](#)

Just implementation of default dynamic array. The same as in C++, not including some Rust-related differences. Asymptotics are the same as you might expect.

Some basic algorithms. **They work on slices too:**

- `sort`, uses modified timsort,  $O(N \log N)$ .
- `sort_unstable`, pdqsort<sup>4</sup>, worst-case  $O(N \log N)$  and best  $O(N)$ .

---

<sup>4</sup>Danila Kutenin blog (Russian). [On pqsort, New sorting algorithm for LLVM libcxx](#)

<sup>5</sup>[GoingNative 2013 C++ Seasoning](#)

<sup>6</sup>[Vec documentation](#)

Just implementation of default dynamic array. The same as in C++, not including some Rust-related differences. Asymptotics are the same as you might expect.

Some basic algorithms. **They work on slices too:**

- `sort`, uses modified timsort,  $O(N \log N)$ .
- `sort_unstable`, pdqsort<sup>4</sup>, worst-case  $O(N \log N)$  and best  $O(N)$ .
- `binary_search`.

---

<sup>4</sup>Danila Kutenin blog (Russian). [On pqsort, New sorting algorithm for LLVM libcxx](#)

<sup>5</sup>[GoingNative 2013 C++ Seasoning](#)

<sup>6</sup>[Vec documentation](#)

Just implementation of default dynamic array. The same as in C++, not including some Rust-related differences. Asymptotics are the same as you might expect.

Some basic algorithms. **They work on slices too:**

- `sort`, uses modified timsort,  $O(N \log N)$ .
- `sort_unstable`, pdqsort<sup>4</sup>, worst-case  $O(N \log N)$  and best  $O(N)$ .
- `binary_search`.
- `select_nth_unstable`, quick select with pdqsort, worst-case  $O(N)$ .

---

<sup>4</sup>Danila Kutenin blog (Russian). [On pqsort, New sorting algorithm for LLVM libcxx](#)

<sup>5</sup>[GoingNative 2013 C++ Seasoning](#)

<sup>6</sup>[Vec documentation](#)

Just implementation of default dynamic array. The same as in C++, not including some Rust-related differences. Asymptotics are the same as you might expect.

Some basic algorithms. **They work on slices too:**

- `sort`, uses modified timsort,  $O(N \log N)$ .
- `sort_unstable`, pdqsort<sup>4</sup>, worst-case  $O(N \log N)$  and best  $O(N)$ .
- `binary_search`.
- `select_nth_unstable`, quick select with pdqsort, worst-case  $O(N)$ .
- There's `_by` and `_by_key` variants of functions to customize comparator.

---

<sup>4</sup>Danila Kutenin blog (Russian). [On pqsort, New sorting algorithm for LLVM libcxx](#)

<sup>5</sup>[GoingNative 2013 C++ Seasoning](#)

<sup>6</sup>[Vec documentation](#)

Just implementation of default dynamic array. The same as in C++, not including some Rust-related differences. Asymptotics are the same as you might expect.

Some basic algorithms. **They work on slices too:**

- `sort`, uses modified timsort,  $O(N \log N)$ .
- `sort_unstable`, pdqsort<sup>4</sup>, worst-case  $O(N \log N)$  and best  $O(N)$ .
- `binary_search`.
- `select_nth_unstable`, quick select with pdqsort, worst-case  $O(N)$ .
- There's `_by` and `_by_key` variants of functions to customize comparator.
- Legendary<sup>5</sup> `rotate_left` and `rotate_right`.

---

<sup>4</sup>Danila Kutenin blog (Russian). [On pqsort, New sorting algorithm for LLVM libcxx](#)

<sup>5</sup>[GoingNative 2013 C++ Seasoning](#)

<sup>6</sup>[Vec documentation](#)

Inside, it's just simple circular deque. Asymptotics are the same as you might expect. **Is not the same** as in C++.

---

<sup>7</sup>[VecDeque documentation](#)

Inside, it's just simple circular deque. Asymptotics are the same as you might expect. **Is not the same** as in C++.

Properties:

---

<sup>7</sup>[VecDeque documentation](#)



Inside, it's just simple circular deque. Asymptotics are the same as you might expect. **Is not the same** as in C++.

Properties:

- The same functions as in **Vec**.

---

<sup>7</sup>[VecDeque documentation](#)

Inside, it's just simple circular deque. Asymptotics are the same as you might expect. **Is not the same** as in C++.

Properties:

- The same functions as in **Vec**.
- And the same core differences from unsafe languages.

---

<sup>7</sup>[VecDeque documentation](#)

Inside, it's just simple circular deque. Asymptotics are the same as you might expect. **Is not the same** as in C++.

Properties:

- The same functions as in `Vec`.
- And the same core differences from unsafe languages.
- Since it's simply circular deque, we have `make_contiguous` and `as_slices`. Moreover, we can easily use optimizations that rely on contiguous element storing, such as SIMD.

---

<sup>7</sup>[VecDeque documentation](#)

## Do you remember `std::deque`?

In C++ standard, we require from `std::deque` not to invalidate iterators when modifying it. This leads to monstrous implementation.

## Do you remember `std::deque`?

In C++ standard, we require from `std::deque` not to invalidate iterators when modifying it. This leads to monstrous implementation.

- **Problem:** We cannot use default circular deque since it invalidates iterators when reallocating.

## Do you remember `std::deque`?

In C++ standard, we require from `std::deque` not to invalidate iterators when modifying it. This leads to monstrous implementation.

- **Problem:** We cannot use default circular deque since it invalidates iterators when reallocating.
- **Solution:** use circular deque, but instead of storing actual element we will store pointer to it.

## Do you remember `std::deque`?

In C++ standard, we require from `std::deque` not to invalidate iterators when modifying it. This leads to monstrous implementation.

- **Problem:** We cannot use default circular deque since it invalidates iterators when reallocating.
- **Solution:** use circular deque, but instead of storing actual element we will store pointer to it.
- **Problem:** we are not cache local, the deque will be really slow.

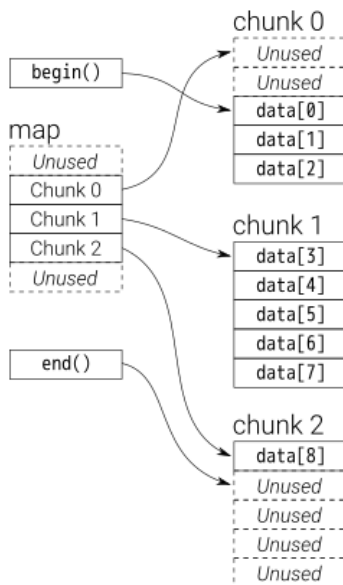
## Do you remember `std::deque`?

In C++ standard, we require from `std::deque` not to invalidate iterators when modifying it. This leads to monstrous implementation.

- **Problem:** We cannot use default circular deque since it invalidates iterators when reallocating.
- **Solution:** use circular deque, but instead of storing actual element we will store pointer to it.
- **Problem:** we are not cache local, the deque will be really slow.
- **Solution:** store chunks of elements in circular deque.



## Do you remember `std::deque`?



## Do you remember `std::deque`?

Although it's quite complex, it works not extremely slow in practice. It *may be* good price for having non-invalidating iterators.<sup>8</sup>

But don't forget: you won't be able to rely on contiguous element storing like in Rust!

---

<sup>8</sup>C++ benchmark - `std::vector` VS `std::list` VS `std::deque`

Rust also includes collections sorted by key: map and set.

---

<sup>9</sup> [BTreeMap documentation](#)

<sup>10</sup> [BTreeSet documentation](#)

Rust also includes collections sorted by key: map and set.

- There's a B-tree data structure inside. Thus it is cache-local and works fast on modern CPUs. Asymptotics for most operations are  $O(\log_B N)$ .

---

<sup>9</sup>[BTreeMap documentation](#)

<sup>10</sup>[BTreeSet documentation](#)

## BTreeMap<sup>9</sup> and BTreeSet<sup>10</sup>

Rust also includes collections sorted by key: map and set.

- There's a B-tree data structure inside. Thus it is cache-local and works fast on modern CPUs. Asymptotics for most operations are  $O(\log_B N)$ .
- Lecturer's humble opinion: these names are much better than `std::map` or `std::set` since they show what data structure it uses. B-tree, hash table, cartesian tree are also maps and sets!

---

<sup>9</sup>[BTreeMap documentation](#)

<sup>10</sup>[BTreeSet documentation](#)

Rust also includes collections sorted by key: map and set.

- There's a B-tree data structure inside. Thus it is cache-local and works fast on modern CPUs. Asymptotics for most operations are  $O(\log_B N)$ .
- Lecturer's humble opinion: these names are much better than `std::map` or `std::set` since they show what data structure it uses. B-tree, hash table, cartesian tree are also maps and sets!
- It is a logic error for a key to be modified in such a way that the key's ordering relative to any other key changes while it is in the map. The behavior resulting from such a logic error is not specified but will not be undefined behavior.

---

<sup>9</sup>[BTreeMap documentation](#)

<sup>10</sup>[BTreeSet documentation](#)

## Do you remember `std::map` and `std::set`?

In C++, we also have ordered map and set. Usually they are implemented using red-black trees.

## Do you remember `std::map` and `std::set`?

In C++, we also have ordered map and set. Usually they are implemented using red-black trees.

- If B-tree data structure is really cool, why don't we use it in C++ standard library, downgrading to red-black trees?



## Do you remember `std::map` and `std::set`?

In C++, we also have ordered map and set. Usually they are implemented using red-black trees.

- If B-tree data structure is really cool, why don't we use it in C++ standard library, downgrading to red-black trees?
- *Once again: iterator invalidation.*

## Do you remember `std::map` and `std::set`?

In C++, we also have ordered map and set. Usually they are implemented using red-black trees.

- If B-tree data structure is really cool, why don't we use it in C++ standard library, downgrading to red-black trees?
- *Once again: iterator invalidation.*
- Since B-tree stores elements in chunks and have smaller depth, it way faster than regular BST.

## Do you remember `std::map` and `std::set`?

In C++, we also have ordered map and set. Usually they are implemented using red-black trees.

- If B-tree data structure is really cool, why don't we use it in C++ standard library, downgrading to red-black trees?
- *Once again: iterator invalidation.*
- Since B-tree stores elements in chunks and have smaller depth, it way faster than regular BST.
- This really hurts! C++ don't have standard fast ordered map and set!

What a language without hash table? Rust have two: `HashMap` and `HashSet`.  
Asymptotics are predictable.

---

<sup>11</sup>CppCon 2017: Matt Kulukundis “Designing a Fast, Efficient, Cache-friendly Hash Table, Step by Step”

<sup>12</sup>[HashMap documentation](#)

<sup>13</sup>[HashSet documentation](#)

## HashMap<sup>12</sup> and HashSet<sup>13</sup>

What a language without hash table? Rust have two: **HashMap** and **HashSet**. Asymptotics are predictable.

- This hash table is quite literally the fastest universal hash table in the world currently existing. It uses quadratic probing and SIMD lookup inside.

---

<sup>11</sup>CppCon 2017: Matt Kulukundis “Designing a Fast, Efficient, Cache-friendly Hash Table, Step by Step”

<sup>12</sup>[HashMap documentation](#)

<sup>13</sup>[HashSet documentation](#)

## HashMap<sup>12</sup> and HashSet<sup>13</sup>

What a language without hash table? Rust have two: **HashMap** and **HashSet**. Asymptotics are predictable.

- This hash table is quite literally the fastest universal hash table in the world currently existing. It uses quadratic probing and SIMD lookup inside.
- More specifically, it's Rust port called Hashbrown of Google SwissTable written in C++. If you're interested in the algorithm, you can watch CppCon talk.<sup>11</sup>

---

<sup>11</sup>CppCon 2017: Matt Kulukundis “Designing a Fast, Efficient, Cache-friendly Hash Table, Step by Step”

<sup>12</sup>[HashMap documentation](#)

<sup>13</sup>[HashSet documentation](#)

What a language without hash table? Rust have two: **HashMap** and **HashSet**. Asymptotics are predictable.

- This hash table is quite literally the fastest universal hash table in the world currently existing. It uses quadratic probing and SIMD lookup inside.
- More specifically, it's Rust port called Hashbrown of Google SwissTable written in C++. If you're interested in the algorithm, you can watch CppCon talk.<sup>11</sup>
- **It is a logic error for a key to be modified in such a way that the key's hash or its equality changes while it is in the map.** The behavior resulting from such a logic error **is not specified**, but **will not result in undefined behavior**.

---

<sup>11</sup>CppCon 2017: Matt Kulukundis “Designing a Fast, Efficient, Cache-friendly Hash Table, Step by Step”

<sup>12</sup>[HashMap documentation](#)

<sup>13</sup>[HashSet documentation](#)

## Do you remember `std::unordered_map` and `std::unordered_set`?

In C++, we also have unordered map and set. Usually they are implemented in some strange way.



## Do you remember `std::unordered_map` and `std::unordered_set`?

In C++, we also have unordered map and set. Usually they are implemented in some strange way.

- Again: if Google Swiss table is really cool, why don't we use it in C++ standard library, downgrading to some strange implementation?

## Do you remember `std::unordered_map` and `std::unordered_set`?

In C++, we also have unordered map and set. Usually they are implemented in some strange way.

- Again: if Google Swiss table is really cool, why don't we use it in C++ standard library, donwgrading to some strange implementation?
- *And once more: iterator invalidation!* At least it prevents us to use open addressing.

## Do you remember `std::unordered_map` and `std::unordered_set`?

In C++, we also have unordered map and set. Usually they are implemented in some strange way.

- Again: if Google Swiss table is really cool, why don't we use it in C++ standard library, donwgrading to some strange implementation?
- *And once more: iterator invalidation!* At least it prevents us to use open addressing.
- This leads to not suitable for production standard library hash tables.

A priority queue implemented with a binary heap. **This will be a max-heap.** The same as in C++. Asymptotics are predictable.

---

<sup>14</sup>[BinaryHeap documentation](#)

A priority queue implemented with a binary heap. **This will be a max-heap.** The same as in C++. Asymptotics are predictable.

- Lecturer's humble opinion: this name is much better than `std::priority_queue` since it shows what data structure it uses. For instance, binomial or fibonacci heaps are also priority queues!

---

<sup>14</sup>[BinaryHeap documentation](#)

A priority queue implemented with a binary heap. **This will be a max-heap.** The same as in C++. Asymptotics are predictable.

- Lecturer's humble opinion: this name is much better than `std::priority_queue` since it shows what data structure it uses. For instance, binomial or fibonacci heaps are also priority queues!
- **It is a logic error for an item to be modified in such a way that the item's ordering relative to any other item changes while it is in a heap.** The behavior resulting from such a logic error **is not specified but will not be undefined behavior.**

---

<sup>14</sup>[BinaryHeap documentation](#)

A doubly-linked list with owned nodes. The same as in C++. Asymptotics are predictable.

- You don't need it in almost any situation. It's slow and not memory efficient. Trust me.
- Since it's not C++, we don't have iterators pointing to elements. It's not convenient.
- Writing your list without unsafe is possible, but quite a challenge! Do it if you want to have a borrow checker as your best friend.<sup>15</sup>

---

<sup>15</sup>[Learn Rust With Entirely Too Many Linked Lists](#)

<sup>16</sup>[LinkedList](#) documentation

# String and &str



A Rust way to store a string. **Totally differs** from C++ `std::string`.

---

<sup>17</sup>[String documentation](#)

A Rust way to store a string. **Totally differs** from C++ `std::string`.

- It's UTF-8-encoded.

---

<sup>17</sup>[String documentation](#)

A Rust way to store a string. **Totally differs** from C++ `std::string`.

- It's **UTF-8-encoded**.
- Growable like a **Vec**. It also made up of three components: a pointer to some bytes, a length, and a capacity. This even gives us many functions same to **Vec**.

---

<sup>17</sup>[String documentation](#)

A Rust way to store a string. **Totally differs** from C++ `std::string`.

- It's **UTF-8-encoded**.
- Growable like a **Vec**. It also made up of three components: a pointer to some bytes, a length, and a capacity. This even gives us many functions same to **Vec**.
- UTF-8 is a variable-width character encoding, so you cannot index it since it's UTF-8. To find N-th symbol, you should iterate over string, parsing code points.

---

<sup>17</sup>[String documentation](#)

## String API

```
struct String {  
    vec: Vec<u8>,  
}  
  
impl String {  
    fn new() -> String;  
    fn with_capacity(capacity: usize) -> String;  
    fn from_utf8(vec: Vec<u8>) -> Result<String, FromUtf8Error>;  
    fn from_utf16(v: &[u16]) -> Result<String, FromUtf16Error>;  
    fn into_bytes(self) -> Vec<u8>;  
    fn as_bytes(&self) -> &[u8];  
}
```

## String in depth

What will this code print?

```
let s = String::from("привет");  
println!("{}", s.len());
```

## String in depth

What will this code print?

```
let s = String::from("привет");  
println!("{}", s.len());
```

This outputs 12, since `.len()` gives count of *bytes* in string.

## String in depth

To work directly with a string like in C++, you must convert it to `Vec<char>`

```
let s = String::from("привет");  
let t = s.chars().collect::<Vec<_>>();  
println!("{:?}", t); // ['п', 'р', 'и', 'в', 'е', 'т']
```

`.chars()` is function that creates iterator over chars of the string.



# The `char` type

The `char` type represents a single character.

More specifically, since “character” isn’t a well-defined concept in Unicode, `char` is a “**Unicode scalar value**”, which is similar to, but not the same as, a “**Unicode code point**”.

```
let mut chars = "é".chars();  
// U+00e9: 'latin small letter e with acute'  
assert_eq!(Some('\u{00e9}'), chars.next());  
assert_eq!(None, chars.next());
```

```
let mut chars = "é".chars();  
// U+0065: 'latin small letter e'  
assert_eq!(Some('\u{0065}'), chars.next());  
// U+0301: 'combining acute accent'  
assert_eq!(Some('\u{0301}'), chars.next());  
assert_eq!(None, chars.next());
```

## char type

The size of `char` is always 4 bytes:

```
assert_eq!(std::mem::size_of::<char>(), 4);
```

## &str

`&str` is a slice type of `String`, similar to `std::string_view`. Just like:

```
let vec = vec![1, 2, 3, 4];  
let vec_slice = &vec[1..3]; // &[2, 3]  
let s = String::from("hello");  
let s_slice = &s[1..3]; // "el"
```

Ok, let's take a UTF-8 slice!

```
let s = String::from("привет");  
let s_slice = &s[1..3];
```

Ok, let's take a UTF-8 slice!

```
let s = String::from("привет");  
let s_slice = &s[1..3];  
// thread 'main' panicked at 'byte index 1 is  
// not a char boundary; it is inside 'п'  
// (bytes 0..2) of `привет`'
```

Ok, let's take a UTF-8 slice!

```
let s = String::from("привет");  
let s_slice = &s[1..3];  
// thread 'main' panicked at 'byte index 1 is  
// not a char boundary; it is inside 'п'  
// (bytes 0..2) of `привет`'
```

That means `&str` also have a UTF-8 invariant checked at runtime.

As a string slice, &str have most functions String have:

```
fn as_bytes(&self) -> &[u8];  
fn chars(&self) -> Chars<'_>;  
fn trim(&self) -> &str;  
fn split<'a, P>(&'a self, pat: P) -> Split<'a, P>;  
fn replace<'a, P>(&'a self, from: P, to: &str) -> String;  
// And so on
```



All string constants are &str.

```
let s: &str = "Hello world!";  
let t1 = s.to_string();  
let t2 = s.to_owned(); // The same as t1
```

# Box and Rc

We are already familiar with `Box` type. Let's check one advanced function:

```
fn leak<'a>(b: Box<T, A>) -> &'a mut T;  
fn into_raw(b: Box<T, A>) -> *mut T;
```

Example:

```
let x = Box::new(41);  
let static_ref: &'static mut usize = Box::leak(x);  
*static_ref += 1;  
assert_eq!(*static_ref, 42);
```

**But stop!** Rust is the safe language, no memory unsafety, no undefined behavior, what's wrong!?

**But stop!** Rust is the safe language, no memory unsafety, no undefined behavior, what's wrong!?

*In reality, when you're creating global objects or interacting with other languages, you **have to** leak objects. Moreover, it's **safe** to leak memory, just not good!*

Rc is single-threaded reference-counting pointer. “Rc” stands for “Reference Counted”.

```
let rc = Rc::new(());  
let rc2 = rc.clone(); // Clones Rc, not what inside!  
let rc3 = Rc::clone(&rc); // The same
```

Rc is dropped when all instances of Rc are dropped.

Primary functions:

```
fn get_mut(this: &mut Rc<T>) -> Option<&mut T>;  
fn downgrade(this: &Rc<T>) -> Weak<T>;  
fn weak_count(this: &Rc<T>) -> usize;  
fn strong_count(this: &Rc<T>) -> usize;
```

References to the variable inside Rc are controlled at runtime:

```
let mut rc = Rc::new(42);  
println!("{}", *rc);  
  
*Rc::get_mut(&mut rc).unwrap() -= 41;  
println!("{}", *rc);  
  
let mut rc1 = rc.clone();  
println!("{}", *rc1);  
// thread 'main' panicked at 'called `Option::unwrap()`  
// on a `None` value'  
// *Rc::get_mut(&mut rc1).unwrap() -= 1;
```

`get_mut` guarantees that it will return mutable reference only if there's only one pointer. If there are more, you won't have a chance to modify Rc.

# Weak

Rc is a **strong** pointer, while Weak is a **weak** pointer. Both of them have *ownership over allocation*, but only Rc have *ownership over the value inside*:

You can upgrade Weak to Rc:

```
fn upgrade(&self) -> Option<Rc<T>>;
```



```
let rc1 = Rc::new(String::from("string"));
let rc2 = rc1.clone();
let weak1 = Rc::downgrade(&rc1);
let weak2 = Rc::downgrade(&rc1);
drop(rc1); // The string is not deallocated
assert!(weak1.upgrade().is_some());
drop(weak1); // Nothing happens
drop(rc2); // The string is deallocated
assert_eq!(weak2.strong_count(), 0);
// If no strong pointers remain, this will return zero.
assert_eq!(weak2.weak_count(), 0);
assert!(weak2.upgrade().is_none());
drop(weak2); // The Rc is deallocated
```

There's also **Arc** - a thread-safe reference-counting pointer. **Arc** stands for "Atomically Reference Counted".

We'll need it to share data safely across threads in the future.

# Conclusion

- We studied API of Option and Result.
- Get acquainted with default collections.
- Learned about smart pointers in Rust.